

Class 06: R Functions

A15827934

All about functions in R

Functions are the way we get stuff done in R. We call a function to read data, compute, plot stuff, etc.

R makes writing functions accessible we should always start by trying to get a working snippet of code first before we write out function.

Todays Lab

We will grade a whole class of student assignments. We will always try to start with a simplified version of the problem.

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)

mean(student1)
```

```
[1] 98.75
```

If we want the average we can use the `mean()` function

```
mean(student1)
```

```
[1] 98.75
```

Let's be super nice instructors and drop the lowest score so the answer here should be 100.

I can use the `min()` function to find the lowest value

```
min(student1)
```

```
[1] 90
```

I found the `which.min()` function that may be useful here. How does it work? Let's just try it:

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
which.min(student1)
```

```
[1] 8
```

```
student1[8]
```

```
[1] 90
```

```
student1[which.min(student1)]
```

```
[1] 90
```

```
student1[1:7]
```

```
[1] 100 100 100 100 100 100 100
```

```
##I can use the minus syntax trick to get  
##everything but the element with the min value.  
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

```
##you cannot do [-8] [1:7] because  
##every student may not have a lowest score  
##at the very end if that makes sense.  
##So good idea to use which.min always as a placeholder
```

```
student1[-which.min(student1)]
```

```
[1] 100 100 100 100 100 100 100
```

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

```
## I have my first working snippet of code
```

Let's test on the other students

```
student2
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
mean(student2[-which.min(student2)])
```

```
[1] NA
```

This gives you N/A. where is the problem??

```
mean(student2)
```

```
[1] NA
```

```
mean(student2, na.rm=TRUE)
```

```
[1] 91
```

oh it is the `mean()` with NA input returns NA by default but I can change

```
student3
```

```
[1] 90 NA NA NA NA NA NA NA
```

```
mean(student3)
```

```
[1] NA
```

```
mean(student3, na.rm=TRUE)
```

```
[1] 90
```

NO bueno. We need to fix this.

I want to stop working with `student1`, `student2` etc and typing it out every time so let instead work with an input called `x`

```
x<-student2  
x
```

```
[1] 100 NA 90 90 90 90 97 80
```

We want to overwrite the NA values with zero - if you miss a hw you score zero on this hw.

ChatGPT and Claude told me about the `is.na()` function. Lets see how it works

```
x
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
is.na(x)
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

```
x[is.na(x)]
```

```
[1] NA
```

We can use logicals to index a vector

```
y<- 1:5  
y
```

```
[1] 1 2 3 4 5
```

```
y>3
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE
```

```
y[y>3]
```

```
[1] 4 5
```

```
y[y>3]<- 100
```

```
x[is.na(x)]<- 0  
x
```

```
[1] 100  0  90  90  90  90  97  80
```

```
mean(x[-which.min(x)])
```

```
[1] 91
```

```
mean(x)
```

```
[1] 79.625
```

```
## it seems that mean(x[-which.min(x)])  
##is the correct one that we want as  
##it ignores the lowest score  
##(which in this case is now zero)  
## We want to drop lowest score and get the mean
```

```
a<-student1  
a
```

```
[1] 100 100 100 100 100 100 100 90
```

```
is.na(a)
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
a[is.na(a)]<- 0  
a
```

```
[1] 100 100 100 100 100 100 100 90
```

```
mean(a[-which.min(a)])
```

```
[1] 100
```

```
b<-student3  
b
```

```
[1] 90 NA NA NA NA NA NA NA
```

```
is.na(b)
```

```
[1] FALSE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
b[is.na(b)]<-0
```

```
#Mask NA values to zero  
x<-student3  
x[is.na(x)]<- 0  
x
```

```
[1] 90 0 0 0 0 0 0 0
```

```
#Drop lowest score and get the mean
mean(x[-which.min(x)])
```

```
[1] 12.85714
```

This is my working snippet of code that solves the problem for all my example student inputs

Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped. Your final function should be adequately explained with code comments and be able to work on an example class gradebook such as this one in CSV format: “<https://tinyurl.com/gradeinput>” [3pts]

```
grade<-function(x){
  #Mask NA values to zero
  x[is.na(x)]<- 0

  #Drop lowest score and get the mean
  mean(x[-which.min(x)])
}
```

Use this function:

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

We need to read the gradebook

```
gradebook<-read.csv("https://tinyurl.com/gradeinput", row.names=1)
gradebook
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	NA	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	NA
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	NA	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

I can use the `apply()` function if I figure out how to use the thing

```
finalgrade<-apply(gradebook,1, grade)
finalgrade
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

```
apply(gradebook,2,grade)
```


hw1	hw2	hw3	hw4	hw5
89.36842	76.63158	81.21053	89.63158	83.42105

Q2. Using your `grade()` function and the supplied gradebook, Who is the top scoring student overall in the gradebook?

```
which.max(finalgrade)
```

```
student-18
18
```

```
##here you would use the apply name not the function
##name grade as function name grade is inside apply
```

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)?

We could calculate the `mean()` score for each homework

```
apply(gradebook,2,mean)
```

hw1	hw2	hw3	hw4	hw5
89.0	NA	80.8	NA	NA

```
apply(gradebook,2,grade)
```

hw1	hw2	hw3	hw4	hw5
89.36842	76.63158	81.21053	89.63158	83.42105

Weird there are some NA. When you do mean. When you do grade function it says HW2

```
##we gonna make into name file because
##we do not want to mess with the original gradebook
mask<-gradebook

mask[is.na(mask)]<- 0
hw.avg<-apply(mask,2,mean)
hw.avg
```

```
hw1 hw2 hw3 hw4 hw5
89.00 72.80 80.80 85.15 79.25
```

```
which.min(hw.avg)
```

```
hw2
2
```

```
apply(gradebook,2,mean,na.rm=T)
```

```
hw1 hw2 hw3 hw4 hw5
89.00000 80.88889 80.80000 89.63158 83.42105
```

Different answer interesting

We could use sum

```
apply(gradebook,2,sum, na.rm=T)
```

```
hw1 hw2 hw3 hw4 hw5
1780 1456 1616 1703 1585
```

Either 2 or 3 depending on what you did. So when you turned NA into 0 you have lower average than when you just ignore or remove NA. The latter would have higher average.

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)?

```
mask$hw2
```

```
[1] 73 64 69 0 100 78 100 100 100 72 66 70 100 100 65 100 63 0 68
[20] 68
```

```
cor(mask$hw2, finalgrade)
```

```
[1] 0.176778
```

```
cor(mask$hw5, finalgrade)
```

```
[1] 0.6325982
```

```
cor(mask$hw1, finalgrade)
```

```
[1] 0.4250204
```

```
cor(mask$hw3, finalgrade)
```

```
[1] 0.3042561
```

```
cor(mask$hw4, finalgrade)
```

```
[1] 0.3810884
```

```
apply(mask,2,cor, y=finalgrade)
```

	hw1	hw2	hw3	hw4	hw5
	0.4250204	0.1767780	0.3042561	0.3810884	0.6325982

```
which.max(apply(mask,2,cor, y=finalgrade)  
)
```

```
hw5  
5
```